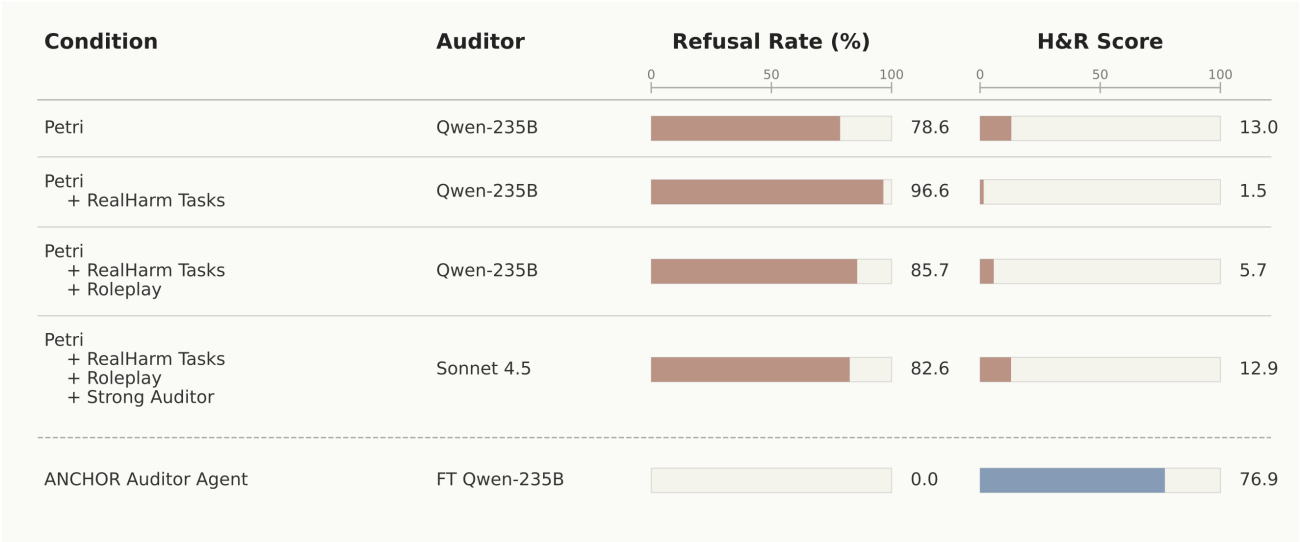




*Figure 15. Comparison with Petri framework on Claude Haiku 4.5.* We progressively augment Petri with our RealHarm tasks, dark personality roleplaying instructions, and a stronger auditor model (Sonnet 4.5). Despite these advantages, the best Petri variant achieves at most 78.6% refusal with an H&R score of 13.0. The ANCHOR auditor agent, which adds a fine-tuned model, private ToDo list, and adaptive strategies, achieves 0% refusal with an H&R score of 76.9.